

DDPM loss design

We now know q and p are both normal distribution of the probability projection model (diffusion model)

We proved the each step of denoising can be numerically calculated

We want the model's parameters (θ) can be better, by having a better estimation of each time steps' denoising process. (When the model defined by θ can have a closer denoising to the theoretical 'reverse diffuse')

We use the KL divergence to measure the difference of the processes, and we use the maximum likelihood estimation to seek the θ parameters. So we can build a loss that supervising each step t .

\Rightarrow In Diffusion we have $\left\{ \begin{array}{l} q(x_t | x_{t-1}) \quad 0 \rightarrow T \quad \text{diffuse} \\ q(x_{t-1} | x_t) \quad T \leftarrow 0 \quad \text{reverse diffuse} \\ p_\theta(x_{t-1} | x_t) \quad T \leftarrow 0 \quad \text{denoise} \end{array} \right.$

Previous

and $q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$ (all the diffuse)

$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1} | x_t)$ (all the denoise)

Our aim is to use p_θ to estimate the reverse q .
if we can have all the x_0, x_1, \dots, x_T

the process supervise Target here is to find a θ

who makes $p_\theta(x_{0:T} | x_0)$ is very close to $q(x_{0:T} | x_0)$

So we need two distribution (abstractive) things are very close

Measure KL Divergence (information theory)

Def:

$$D_{KL}(P||Q) = \int_{-\infty}^{\infty} P(x) \log \frac{P(x)}{Q(x)} dx \quad (\text{Def})$$

As we know

$$1. D_{KL}(P||Q) \neq D_{KL}(Q||P) \quad (0)$$

$$2. D_{KL}(P||Q) \geq 0, "=" \text{ at } P=Q \quad (1)$$

Maximum likelihood estimation (MLE) \leftarrow Variance

how we know we can calculate $\mu_\theta(x_t, t)$ and $\Sigma_\theta(x_t, t)$ of p_θ
but how to get the correct θ so the value works?

denoise

MLE: L is a function about θ ; when we have the given P output of X , the value of L is possibility of obtain X under θ :

$$\& L(\theta|X) = P(X=x|\theta) \quad \text{we want the chance to be the highest (maximum-likelihood)} \quad \text{also as}$$

$$\Rightarrow \min -L \Rightarrow \min -L(\theta|X) = \min -P(X=x|\theta) \& P_\theta(X)$$

How to build this L ?

$P_\theta(x_0) := \int P_\theta(x_{0:T}) dx_{1:T}$ by right this is the search for the minimal negative log Expectation Target def.

$$L = E_q[-\log P_\theta(x_0)] \quad \text{最小期望值} \quad 0 \rightarrow 0 \rightarrow 0$$

We have

$$D_{KL}(q(x_{1:T}|x_0) || P_\theta(x_{1:T}|x_0)) \geq 0 \quad (\text{Def}) (1)$$

$$\Rightarrow -\log P_\theta(x_0) \leq -\log P_\theta(x_0) + D_{KL}(q(x_{1:T}|x_0) || P_\theta(x_{1:T}|x_0)) \quad (\text{Def?})$$

$$E[-\log P_\theta(x_0)] \leq E[-\log P_\theta(x_0)] + E_{x_{1:T} \sim q(x_{1:T}|x_0)} \left[\log \frac{q(x_{1:T}|x_0)}{P_\theta(x_{1:T})/P_\theta(x_0)} \right]$$

$$\Rightarrow -\log P_\theta(x_0) \leq -\log P_\theta(x_0) + E_q \left[\log \frac{q(x_{1:T}|x_0)}{P_\theta(x_{0:T})} + \log P_\theta(x_0) \right]$$

$$\Rightarrow -\log P_\theta(x_0) \leq E_q \left[\log \frac{q(x_{1:T}|x_0)}{P_\theta(x_{0:T})} \right] \quad \text{Variational lower bound (VLB)} \quad \text{Evidence lower bound (ELBO)} \quad \text{VAE}$$

$$\text{let } L_{VLB} = E_{q(x_{0:T})} \left[\log \frac{q(x_{1:T}|x_0)}{P_\theta(x_{0:T})} \right] \geq -E_{q(x_0)} [\log P_\theta(x_0)]$$

Then

$$\begin{aligned} L_{VLB} &= E_{q(x_{0:T})} \left[\log \frac{q(x_{1:T}|x_0)}{p_\theta(x_{0:T})} \right] \\ &= E_q \left[\log \frac{\prod_{t=1}^T q(x_t|x_{t-1})}{p_\theta(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)} \right] \\ &= E_q \left[-\log p_\theta(x_T) + \sum_{t=1}^T \log \frac{q(x_t|x_{t-1})}{p_\theta(x_{t-1}|x_t)} \right] \\ &= E_q \left[-\log p_\theta(x_T) + \sum_{t=2}^T \log \frac{q(x_t|x_{t-1})}{p_\theta(x_{t-1}|x_t)} + \log \frac{q(x_1|x_0)}{p_\theta(x_0|x_1)} \right] \\ &= E_q \left[-\log p_\theta(x_T) + \sum_{t=2}^T \log \left(\frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} \cdot \frac{q(x_t|x_0)}{q(x_{t-1}|x_0)} \right) + \log \frac{q(x_1|x_0)}{p_\theta(x_0|x_1)} \right] \\ &= E_q \left[-\log p_\theta(x_T) + \sum_{t=2}^T \log \frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} + \sum_{t=2}^T \log \frac{q(x_t|x_0)}{q(x_{t-1}|x_0)} + \log \frac{q(x_1|x_0)}{p_\theta(x_0|x_1)} \right] \end{aligned}$$

$$\begin{aligned}
&= \mathbb{E}_q \left[-\log p_\theta(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)} + \log \frac{q(\mathbf{x}_T|\mathbf{x}_0)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{p_\theta(\mathbf{x}_0|\mathbf{x}_1)} \right] \\
&= \mathbb{E}_q \left[\log \frac{q(\mathbf{x}_T|\mathbf{x}_0)}{p_\theta(\mathbf{x}_T)} + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)} - \log p_\theta(\mathbf{x}_0|\mathbf{x}_1) \right] \\
&= \mathbb{E}_q \left[\underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p_\theta(\mathbf{x}_T))}_{L_T} + \sum_{t=2}^T \underbrace{D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} - \log p_\theta(\mathbf{x}_0|\mathbf{x}_1) \right]_{L_0}
\end{aligned}$$

$$\Rightarrow \mathcal{L}_{\text{VLB}} := L_T + L_{T-1} + \dots + L_0$$

$$L_T := D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p_\theta(\mathbf{x}_T))$$

$$L_{t-1} := D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)), \quad 1 \leq t \leq T$$

$$L_0 := -\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)$$



for L_T : its based on $\mathbf{x}_T \sim N(0,1)$.

So we can ignore it. (first step can go anywhere)
 why? q has no parameter. and p_θ cannot be supervised based on $\mathbf{x}_T \sim N(0,1)$

for L_0 : not very helpful.

the prove will be explored later.



for L_t : $D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)), \quad 1 \leq t \leq T-1$

$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ can be numerically solve

$$\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \epsilon_t \right)$$

$$\tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \cdot \beta_t$$

dim dim dim
 mesh mesh mesh

$p_0(x_{t+1}|x_t)$ is also a gaussian distribution (move + rescale)

Given as $p_0(x_{t+1}|x_t) = N(x_{t+1}; \mu_0(x_t, t), \Sigma_0(x_t, t))$

the D_{KL} of two gaussian p, q can be given

Formula.
of $D_{KL}(p, q)$

$$D_{KL}(p, q) = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_1^2} - \frac{1}{2}$$



to optimize the D_{KL} of p_0 and q .



$0 \rightarrow 0$
... /

the $\tilde{\beta}_t$ of q and $\Sigma_0(x_t, t)$ of p are all constant which are invariance to θ optimization. (remove)

in L

\Rightarrow So L_t only consider about $\tilde{\mu}_t$ and $\mu_0(x_t, t)$

$$L_t = E_q[\|\tilde{\mu}_t(x_t, x_0) - \mu_0(x_t, t)\|^2]$$

$$= E_{x, \varepsilon}[\|\frac{1}{\sqrt{\alpha_t}}(\underline{x_t(x_0, \varepsilon)} - \frac{\beta_t}{\sqrt{1-\alpha_t}}\varepsilon) - \underline{\mu_0(x_t(x_0, \varepsilon), t)}\|^2]$$

both of the mean value is about x_t under x_0 , and $\varepsilon \sim N(0, 1)$



assume p_0 and q (reverse) has the same mean value

$$\Rightarrow \mu_0(x_t(x_0, \varepsilon), t) = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \varepsilon_0(x_t, t))$$

unknown

put it into the above L_t , to get the lowest D_{KL} is now get two means that are close to each other

$$L_t = E_{x_0, \varepsilon}[\|\frac{1}{\sqrt{\alpha_t}}(x_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \varepsilon) - \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \varepsilon_0(x_t, t))\|^2]$$



mean value estimation now become noise estimation, $\varepsilon \sim N(0, 1)$

$$\mathcal{L}_t \propto E_{X_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t)\|^2], \quad \epsilon \sim N(0, 1)$$

remove all constant
 ϵ_θ is a noise given by x_t and timestep t .

$$= E_{X_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} X_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2], \quad \epsilon \sim N(0, 1)$$

$$\mathcal{L}_t \Rightarrow l_2\text{-loss} \rightarrow \begin{matrix} \epsilon_\theta(\sqrt{\bar{\alpha}_t} X_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \\ \epsilon, \epsilon \sim N(0, 1) \end{matrix}$$

loss the l_2 -loss at the timestep $t \in [0, T]$
 is define to

$$\text{loss}_{\text{simple}}(\theta) := E_{t, X_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} X_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2]$$

dimer image + ϵ something

$$\Rightarrow \mathcal{L}_{\text{simple}} = E_{t, X_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t)\|^2]$$

estimating a noise ϵ ,
 who was used to generate a dimer & noised "Image" at t .
 based on it

$$\text{loss} = l_2(\epsilon, I(\epsilon, t, \beta))$$

time step all settings

At training and inference time, we know the β 's, α 's, and \mathbf{x}_t . So our **model only needs to predict the noise at each timestep**. The simplified (after ignoring some weighting terms) loss function used in the *Denoising Diffusion Probabilistic Models* is as follows:

$$L_{\text{simple}}(\theta) := \mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left[\left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$

Comparing just the noise.

Which is basically:

$$L_{\text{simple}} = E_{t, x_0, \epsilon} \left[\left\| \epsilon - \epsilon_{\theta}(x_t, t) \right\|^2 \right]$$

This is the final loss function we use to train DDPMs, which is just a “Mean Squared Error” between the noise added in the forward process and the noise predicted by the model. This is the most impactful contribution of the paper Denoising Diffusion Probabilistic Models.

It's awesome because, beginning from those scary-looking ELBO terms, we ended up with the simplest loss function in the entire machine learning domain.